

Predicting Car Crash and Cost data From Insurance Data

Daniel Brooks (daniel.brooks@spsmail.cuny.edu), Daniel Fanelli (daniel.fanelli@spsmail.cuny.edu), Christopher Fenton (christopher.fenton@spsmail.cuny.edu), James Hamski (james.hamski@spsmail.cuny.edu), Youqing Xiang (youqing.xiang@spsmail.cuny.edu)

7/10/2016

1 Introduction

The purpose of this assignment is to build a series two different models. The first model will predict whether a person will get into a car crash and the second model will be used to predict the amount as to which the crash will cost. Both of these models will utilize the customers insurance information to predict the two variables stated above.

Our data consists of 8161 customers with 24 potential predictor variables and 2 target (response) variables. The response variables are:

TARGET_FLAG: This flag will take a value of 0 or 1. One means that the person was in a car crash and 0 means that they were not.

TARGET_AMT: This is the amount associated with the car crash. If the TARGET_FLAG field is 0, then there will be no amount, because the person did not get into a crash. If the TARGET_FLAG is a 1, then there will be a value in this field, because there was a crash.

2 Data Exploration

The first thing that has to be taken care of is the missing (NA) values in the data. We can see that in Figure 1.1, there are a few fields that have some missing values. We can see that the fields JOB, CAR_AGE, HOME_VAL, YOJ, AGE and INCOME have missing values. These need to either be removed or imputed to continue on with the analysis.

Missing Values vs Observed

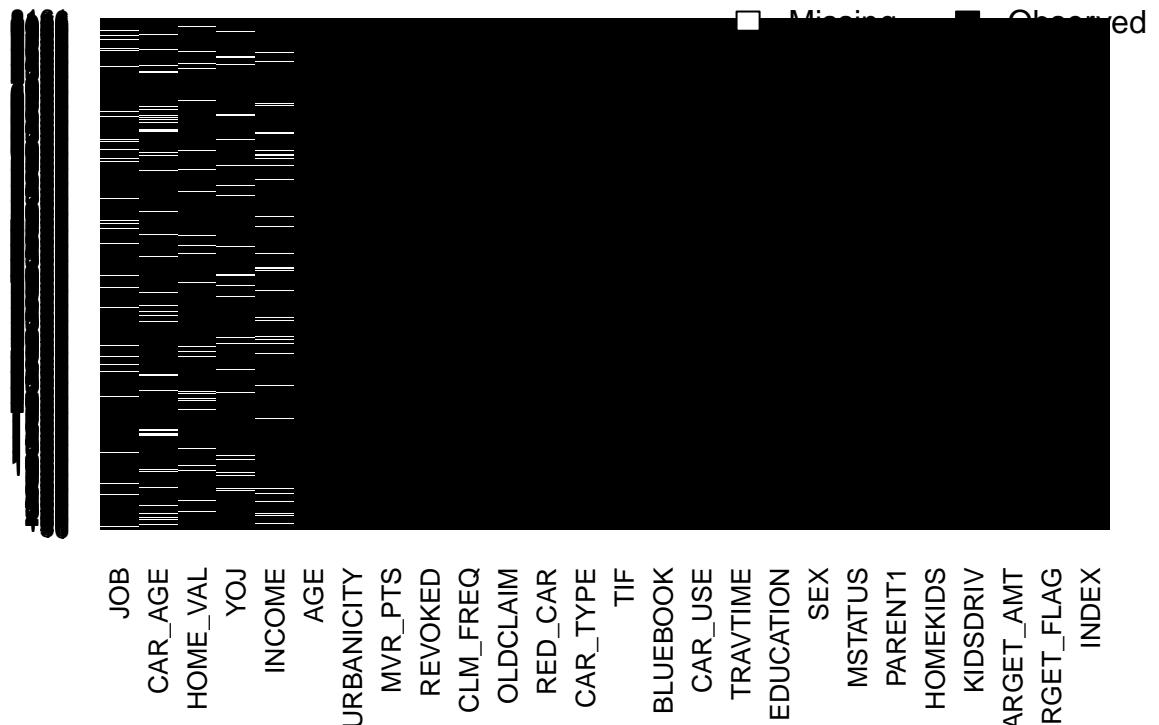


Figure 1.1

We can also see, in the summary statistics (Figure 1.2) some of the data is incorrect. The CAR_AGE field is showing that the minimum car age is -3. That is not a possible value for that field. That field will have to be looked at more closely to see what other values are invalid. They will have to be re-imputed or removed before the final models are created.

The data also contains 10 categorical variables and 16 numeric variables. The categorical data will need to be converted into a numerical field in order to be plugged into the model and get a value for the two target variables.

Figure 1.2

```

##      INDEX      TARGET_FLAG      TARGET_AMT      KIDSDRV
##  Min.   : 1  Min.   :0.0000  Min.   : 0  Min.   :0.0000
##  1st Qu.: 2559  1st Qu.:0.0000  1st Qu.: 0  1st Qu.:0.0000
##  Median : 5133  Median :0.0000  Median : 0  Median :0.0000
##  Mean   : 5152  Mean   :0.2638  Mean   : 1504  Mean   :0.1711
##  3rd Qu.: 7745  3rd Qu.:1.0000  3rd Qu.: 1036  3rd Qu.:0.0000
##  Max.   :10302  Max.   :1.0000  Max.   :107586  Max.   :4.0000
##
##      AGE      HOMEKIDS      YOJ      INCOME
##  Min.   :16.00  Min.   :0.0000  Min.   : 0.0  Length:8161
##  1st Qu.:39.00  1st Qu.:0.0000  1st Qu.: 9.0  Class :character
##  Median :45.00  Median :0.0000  Median :11.0  Mode  :character
##  Mean   :44.79  Mean   :0.7212  Mean   :10.5
##  3rd Qu.:51.00  3rd Qu.:1.0000  3rd Qu.:13.0
##  Max.   :81.00  Max.   :5.0000  Max.   :23.0
##  NA's   : 6      NA's   :454

```

```

##      PARENT1          HOME_VAL        MSTATUS
## Length:8161      Length:8161      Length:8161
## Class :character  Class :character  Class :character
## Mode  :character  Mode  :character  Mode  :character
##
##      SEX            EDUCATION        JOB          TRAVTIME
## Length:8161      Length:8161      Length:8161      Min.   : 5.00
## Class :character  Class :character  Class :character  1st Qu.: 22.00
## Mode  :character  Mode  :character  Mode  :character  Median  : 33.00
##                                         Mean   : 33.49
##                                         3rd Qu.: 44.00
##                                         Max.   :142.00
##
##      CAR_USE         BLUEBOOK        TIF          CAR_TYPE
## Length:8161      Length:8161      Min.   : 1.000  Length:8161
## Class :character  Class :character  1st Qu.: 1.000  Class :character
## Mode  :character  Mode  :character  Median  : 4.000  Mode  :character
##                                         Mean   : 5.351
##                                         3rd Qu.: 7.000
##                                         Max.   :25.000
##
##      RED_CAR         OLDCALLM       CLM_FREQ      REVOKED
## Length:8161      Length:8161      Min.   :0.0000  Length:8161
## Class :character  Class :character  1st Qu.:0.0000  Class :character
## Mode  :character  Mode  :character  Median  :0.0000  Mode  :character
##                                         Mean   :0.7986
##                                         3rd Qu.:2.0000
##                                         Max.   :5.0000
##
##      MVR_PTS         CAR_AGE        URBANICITY
## Min.   : 0.000  Min.   :-3.000  Length:8161
## 1st Qu.: 0.000  1st Qu.: 1.000  Class :character
## Median : 1.000  Median : 8.000  Mode  :character
## Mean   : 1.696  Mean   : 8.328
## 3rd Qu.: 3.000  3rd Qu.:12.000
## Max.   :13.000  Max.   :28.000
## NA's   :510

```

In Figure 1.3 we take a look at the first target variable (TARGET_FLAG). The summary stats are saying that the mean is less than .5 which means that the data is not evenly disbursed between 0 and 1. The mean would be .5 if that were the case. We can see by the histogram that there are about 3 times as many 0's (no crashes) then there are crashes in this data set. In Figure 1.4, we see that the amount that the customer paid is also around 0. That is agreeing with the high amount of no crashes that we saw in Figure 1.3. This means we can proceed with our analysis.

Figure 1.3

```

##      Min. 1st Qu. Median   Mean 3rd Qu.   Max.
## 0.0000 0.0000 0.0000 0.2638 1.0000 1.0000

```

Histogram of data\$TARGET_FLAG

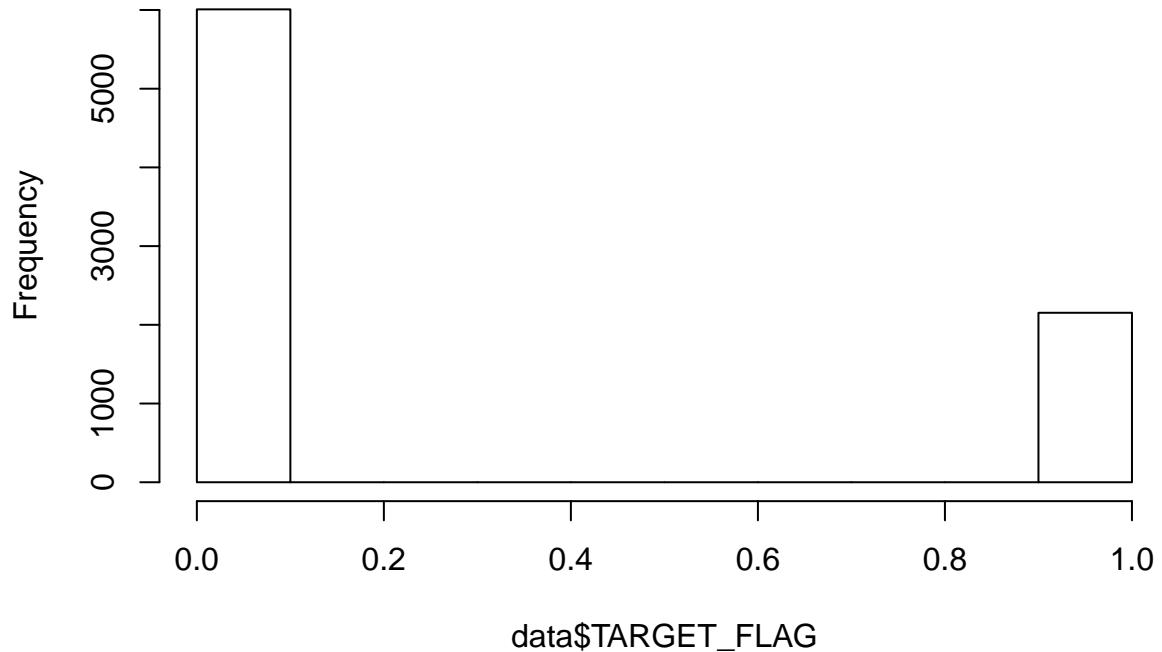
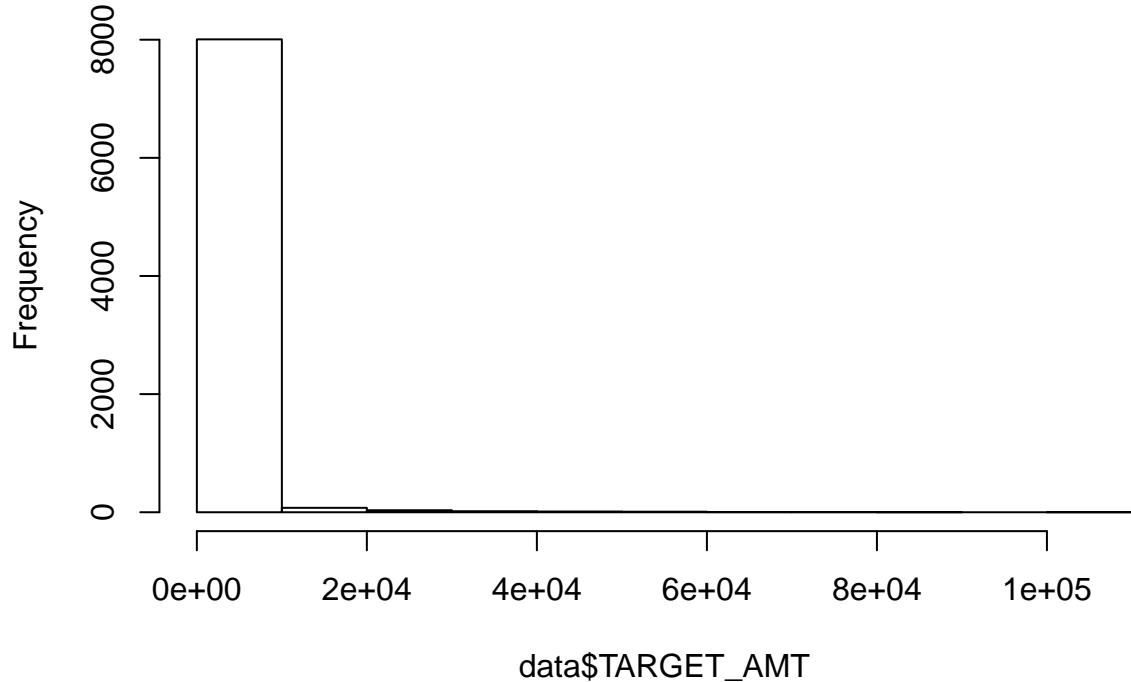


Figure 1.3

```
##      Min. 1st Qu. Median     Mean 3rd Qu.    Max.  
##        0       0       0     1504    1036 107600
```

Histogram of data\$TARGET_AMT



Finally, we want to look at the variables and see how they correlate with each other. In Figure 1.5 we see a grid that plots each field in the data set against each other. It also shows how strong of a correlation there is between each variable through the correlation value and the amount of stars in each box (correlation significance), 3 stars means that the variables have a very strong correlation and the strength drops as the number of stars drop.

We can see from the grid, that not all of the variables correlate with each other. The variables that appear to not make good predictors would be, MVR_PTS, TRAVTIME, and TIF. These variables have a low correlation significance. This means that they may not be good choices to include in our final model for any of the target variables. All of the other variables appear to have a relatively strong correlation between each other and would be good choices for the overall model.

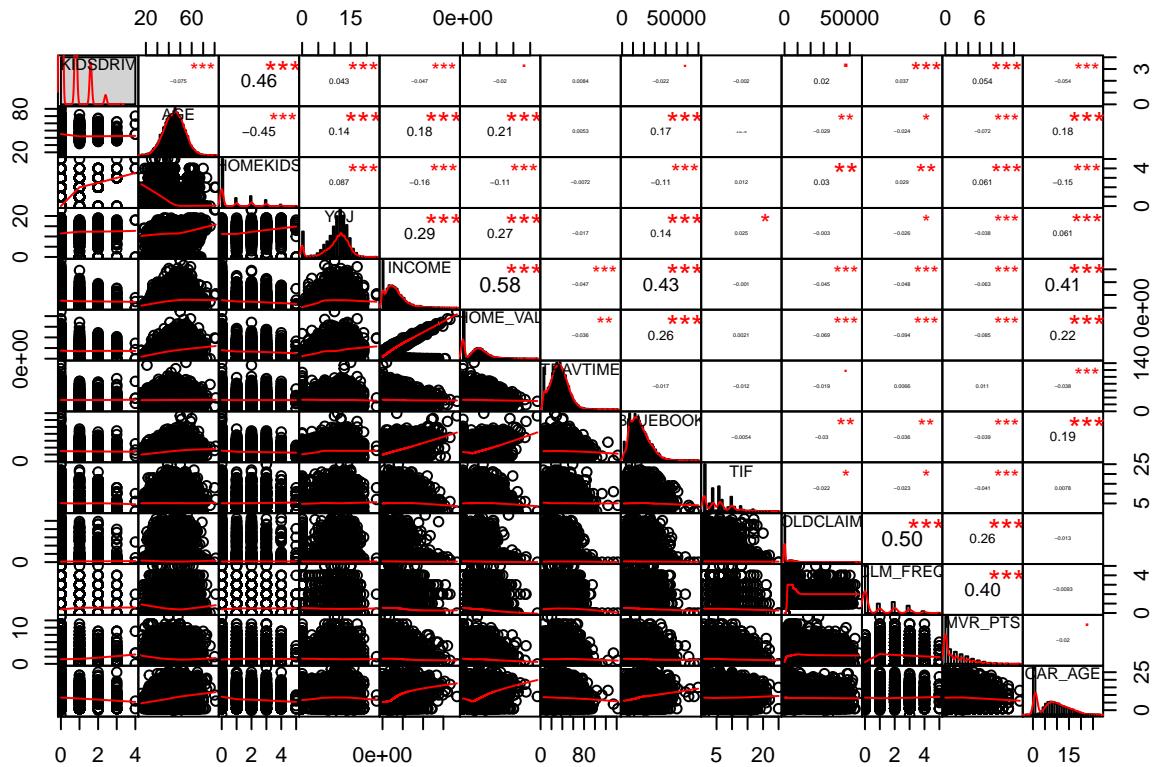


Figure 1.5

Figure 1.6 is another way to show the correlation between variables, just like in Figure 1.5. This grid has different color/size circles within each variable intersection in the grid. The bigger the circle the better the correlation. We can see that the fields TIF, TRAVTIME, and MVR PTS are still relatively bad choices for variables, just like in Figure 1.5.

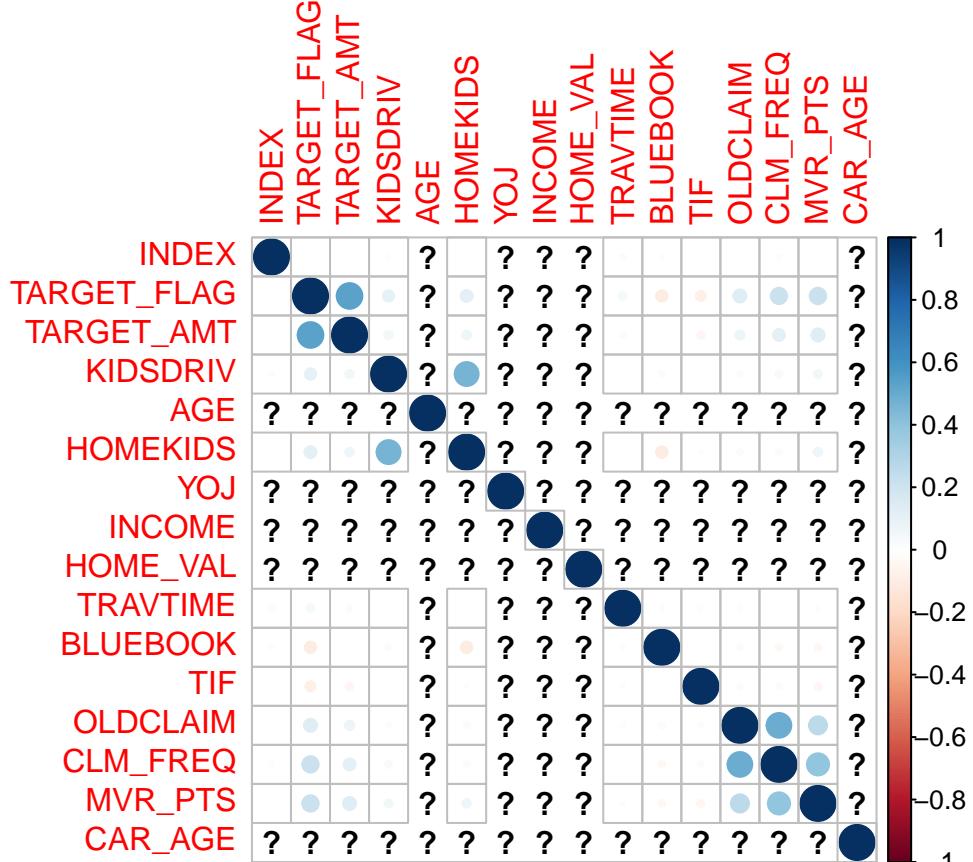


Figure 1.6

3 Data Preparation

3.1 Dealing with NA's and Invalid Values

The first thing that we want to take care of is the NA's (missing values) in the data. We decided that the best course of action would be to remove any of the records that were not complete. Also, we will remove any of the data that is invalid (like a car age below 0). With the removal of the NA's and the invalid values, the total observations for the data is reduced from 8161 to 6044 total observations. That is still plenty of observations to do analysis and modeling on.

3.2 Categorical Data

Within our data, there are 10 fields that are categorical. These variables group the data into different sections. The best way to deal with this data is to make dummy variables. Most of the data contains 1 or 2 possibilities (YES/NO etc). This allows me to assign a value of 0 for one possibility and a 1 for another. It changes the factor data into a number so it can be used in the regression analysis. The education and jobs fields were a little different. They have more than 2 possibilities. I grouped the possibilities into logical groups. The education was grouped by level of academic achievement (Masters and above = 1, below is a 0). The education was grouped along the same way, a college education/advanced education was grouped into one (Lawyer, Professional, Manager = 1), everybody else gets a 0. The field of CAR_TYPE is dealt the same way as well. Those who drive a Panel Truck, Pickup or Sports car will be labeled with a 1 and everything else will be labeled with a 0.

3.3 Accounting Type Data

The last thing that needed to be done with the data is to convert the accounting data, into numeric fields. The accounting data (dollar amount data) has dollar signs within the data. That means that the data will be treated as a character set. The dollar sign needs to be removed and the data needs to be changed into a number, so it can be used in the analysis later on.

4 Build Models

4.1 TARGET_FLAG

This is the target variable that will tell us there was a crash or not for the given customer. If the field takes a value of 0, that means that the customer was not in an accident, or the accident was not their fault. If the field takes a value of 1, that means the customer has been in an accident, or the accident was their fault.

The first thing that we need to do is split the data up into a training set and a test set. We will be taking the data and separating it up so 70% of the data is the training set, and 30% of the data will be the testing set.

4.1.1 Model 1

In this model we are doing a type of stepwise function. We are taking the training set and removing some of the variables that we feel are not good predictors. We run the model against the remaining variables and look at the output. Then, we go back through and remove anymore variables that we do not feel are good predictors. That will leave us with a final function that has all the variables that we fell will make the best predicting function.

The output for this function is as follows:

```
##  
## Call:  
## lm(formula = TARGET_FLAG ~ TARGET_AMT + PARENT1 + HOME_VAL +  
##      MSTATUS + TRAVTIME + CAR_USE + REVOKED + CAR_AGE + URBANICITY,  
##      data = training_2a)  
##  
## Residuals:  
##      Min        1Q    Median        3Q       Max  
## -3.1000 -0.2157 -0.1139  0.1073  0.9780  
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 2.497e-01 2.769e-02  9.018 < 2e-16 ***  
## TARGET_AMT  4.461e-05 1.176e-06 37.938 < 2e-16 ***  
## PARENT1     -1.091e-01 1.792e-02 -6.086 1.26e-09 ***  
## HOME_VAL    -4.705e-07 5.026e-08 -9.360 < 2e-16 ***  
## MSTATUS     -2.001e-02 1.363e-02 -1.468 0.142264  
## TRAVTIME    1.681e-03 3.407e-04  4.934 8.38e-07 ***  
## CAR_USE     -1.189e-01 1.157e-02 -10.278 < 2e-16 ***  
## REVOKED     1.120e-01 1.624e-02  6.898 6.06e-12 ***  
## CAR_AGE     -3.572e-03 1.015e-03 -3.518 0.000439 ***  
## URBANICITY   1.978e-01 1.370e-02 14.433 < 2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

## 
## Residual standard error: 0.3476 on 4221 degrees of freedom
## Multiple R-squared:  0.3754, Adjusted R-squared:  0.3741
## F-statistic: 281.9 on 9 and 4221 DF,  p-value: < 2.2e-16

```

4.1.2 Model 2

In this model, we are utilizing a backwards approach into solving for the overall model. The backwards approach to variable selection starts off with all variables in the model. I then starts to remove fields, until it gets to a point where removing anymore fields will not be beneficial to the model. That is the point when the final model is found.

To solve for the TARGET_FLAG field, I will be using a probit function. This function is very useful when there are only two possible outcomes for the field that you are trying to predict. This model utilizes backwards selection when picking the variables for the model. It starts out all of the variables that are possible. It then starts to remove variables until it reaches the optimal solution for the function. The output from the model is as follows:

```

## 
## Call:
## glm(formula = TARGET_FLAG ~ KIDSDRV + AGE + INCOME + PARENT1 +
##      HOME_VAL + MSTATUS + JOB + TRAVTIME + CAR_USE + BLUEBOOK +
##      TIF + CAR_TYPE + RED_CAR + OLDCLAIM + CLM_FREQ + REVOKED +
##      MVR PTS + URBANICITY, family = binomial(link = "probit"),
##      data = train1)
## 
## Deviance Residuals:
##    Min      1Q  Median      3Q     Max
## -2.2830 -0.7328 -0.4145  0.6418  3.3760
## 
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -6.332e-01  1.670e-01 -3.791 0.000150 ***
## KIDSDRV      1.642e-01  4.366e-02  3.760 0.000170 ***
## AGE         -4.428e-03  2.905e-03 -1.524 0.127448
## INCOME      -3.533e-06  7.734e-07 -4.568 4.93e-06 ***
## PARENT1     -2.815e-01  7.865e-02 -3.579 0.000345 ***
## HOME_VAL    -1.038e-06  2.857e-07 -3.635 0.000278 ***
## MSTATUS      1.652e-01  6.474e-02  2.551 0.010732 *
## JOB          -1.912e-01  4.961e-02 -3.854 0.000116 ***
## TRAVTIME     8.263e-03  1.485e-03  5.564 2.64e-08 ***
## CAR_USE      -5.431e-01  5.275e-02 -10.295 < 2e-16 ***
## BLUEBOOK    -1.838e-05  3.276e-06 -5.612 2.00e-08 ***
## TIF          -3.380e-02  5.799e-03 -5.829 5.57e-09 ***
## CAR_TYPE     1.915e-01  5.003e-02  3.828 0.000129 ***
## RED_CAR      -2.533e-01  5.368e-02 -4.719 2.37e-06 ***
## OLDCLAIM    -5.928e-06  3.152e-06 -1.880 0.060055 .
## CLM_FREQ     1.148e-01  2.350e-02  4.883 1.05e-06 ***
## REVOKED      4.657e-01  7.379e-02  6.310 2.79e-10 ***
## MVR PTS     6.991e-02  1.120e-02  6.244 4.26e-10 ***
## URBANICITY   1.170e+00  7.651e-02 15.294 < 2e-16 ***
## --- 
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
## 
```

```

## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 4859.5 on 4230 degrees of freedom
## Residual deviance: 3807.1 on 4212 degrees of freedom
## AIC: 3845.1
##
## Number of Fisher Scoring iterations: 5

```

4.1.3 Model 3

The model is a little interesting in the fact that, the original analysis of the data showed that the fields TIF, TRAVTIME and MVR_PTS were not good predictors. They all had very few starts when being correlated with the other variables. They all showed up in the model. This could be due to the fact that the original analysis was based on a linear comparision. This equation is a log function. That shows that those variables may have a horrible linear relationship with the rest of the variables but a very stong logistic relationship.

To solve for the TARGET_FLAG field, I will be using a probit function, just like during the backwards function. This function goes the opposite way as the backward function. It starts with a plain function and adds variables until it gets to the optimal solution. Once it cannot add variables to make the equation better, it stops and that is the final function. The output for the forward stepping function is as follows:

```

##
## Call:
## glm(formula = TARGET_FLAG ~ URBANICITY + HOME_VAL + CAR_USE +
##      MVR PTS + BLUEBOOK + PARENT1 + REVOKED + TIF + TRAVTIME +
##      INCOME + RED_CAR + CLM_FREQ + JOB + CAR_TYPE + KIDSDRIV +
##      MSTATUS + OLDCLAIM + AGE, family = binomial(link = "probit"),
##      data = train1)
##
## Deviance Residuals:
##    Min      1Q   Median      3Q     Max
## -2.2830 -0.7328 -0.4145  0.6418  3.3760
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -6.332e-01 1.670e-01 -3.791 0.000150 ***
## URBANICITY  1.170e+00 7.651e-02 15.294 < 2e-16 ***
## HOME_VAL    -1.038e-06 2.857e-07 -3.635 0.000278 ***
## CAR_USE     -5.431e-01 5.275e-02 -10.295 < 2e-16 ***
## MVR PTS    6.991e-02 1.120e-02  6.244 4.26e-10 ***
## BLUEBOOK   -1.838e-05 3.276e-06 -5.612 2.00e-08 ***
## PARENT1    -2.815e-01 7.865e-02 -3.579 0.000345 ***
## REVOKED    4.657e-01 7.379e-02  6.310 2.79e-10 ***
## TIF        -3.380e-02 5.799e-03 -5.829 5.57e-09 ***
## TRAVTIME   8.263e-03 1.485e-03  5.564 2.64e-08 ***
## INCOME     -3.533e-06 7.734e-07 -4.568 4.93e-06 ***
## RED_CAR    -2.533e-01 5.368e-02 -4.719 2.37e-06 ***
## CLM_FREQ   1.148e-01 2.350e-02  4.883 1.05e-06 ***
## JOB        -1.912e-01 4.961e-02 -3.854 0.000116 ***
## CAR_TYPE   1.915e-01 5.003e-02  3.828 0.000129 ***
## KIDSDRIV  1.642e-01 4.366e-02  3.760 0.000170 ***
## MSTATUS    1.652e-01 6.474e-02  2.551 0.010732 *
## OLDCLAIM   -5.928e-06 3.152e-06 -1.880 0.060055 .
## AGE        -4.428e-03 2.905e-03 -1.524 0.127448

```

```

## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 4859.5 on 4230 degrees of freedom
## Residual deviance: 3807.1 on 4212 degrees of freedom
## AIC: 3845.1
##
## Number of Fisher Scoring iterations: 5

```

This function is very interesting. In the backwards function, stated above (Model 1), the MSTATUS (Marital status) field is positive, but it is negative in the forwards function. The sign switched depending on which way you come at the optimal function. Since the model also contains the field of KIDSDRV and PARENT1, it is possible that multicollinearity could be playing a factor. It is also very interesting that the TRAVTIME, TIF and MVR_PTS fields are also included in this model just like in the backwards model.

4.1.4 Model 4 (Remove Redundancy, Calculate All Combinations)

Because the data contains such a large number of categorical fields, it is impossible (with our machine power + time constraints) to calculate and test the model for all possible combinations of the categorical's inclusion and/or values. However, what if some of the fields could *simply be removed*? Let's take a look at the *correlations* between some of the dependent attributes in our data.

col_names	accuracy
CAR_TYPE	0.7387707
CAR_USE	0.7387707
EDUCATION	0.7387707
JOB	0.7387707
MSTATUS	0.7387707
PARENT1	0.7387707
RED_CAR	0.7387707
REVOKE	0.7387707
SEX	0.7387707
URBANICITY	0.7387707

Those categorical variables have the exact same calculated accuracties.

Could this be coincidence? Seems those columns could be duplicates, let's run Chi-Square correlation tests to confirm:

all_chi_sq_labels	all_chi_sq_results
car_type_to_car_use	0.000
car_type_to_education	0.000
car_type_to_job	0.759
car_type_to_mstatus	0.160
car_type_to_parent1	0.911
car_type_to_red_car	0.000
car_type_to_revoked	0.940
car_type_to_sex	0.000
car_type_to_urbancity	0.582

Chi-Squared says some columns are exact duplicates: Remove Them! (We'll keep CARTYPE):

Why can we simply remove fields? Imagine there is a categorical called letters with “a”, “b”, “c”, “d” values. And there is a categorical called numbers with “one”, “two”, “three”, “four”. If for EVERY occurrence:

“a” => “one” “b” => “two” “c” => “three” “d” => “four”

Then would it make sense to calculate a model for all 16 letter-to-number combinations? Of course not, they are 100% correlated. We can just keep one of the fields.

```
drops <- c("CAR_USE", "EDUCATION", "JOB", "PARENT1", "RED_CAR", "SEX", "URBANICITY")
ins.train <- ins.train[ , !(names(ins.train) %in% drops)]
ins.test <- ins.test[ , !(names(ins.test) %in% drops)]
```

Also, *car_type_to_mstatus* and *car_type_to_revoked* were close, lets compare mstatus to revoked:

```
mstatus_to_revoked <- chisq.test(table(ins.train$MSTATUS, ins.train$REVOKE))
round(mstatus_to_revoked$p.value, digits = 3)
```

```
## [1] 0.115
```

So those are ALSO **DUPPLICATES**, remove REVOKE

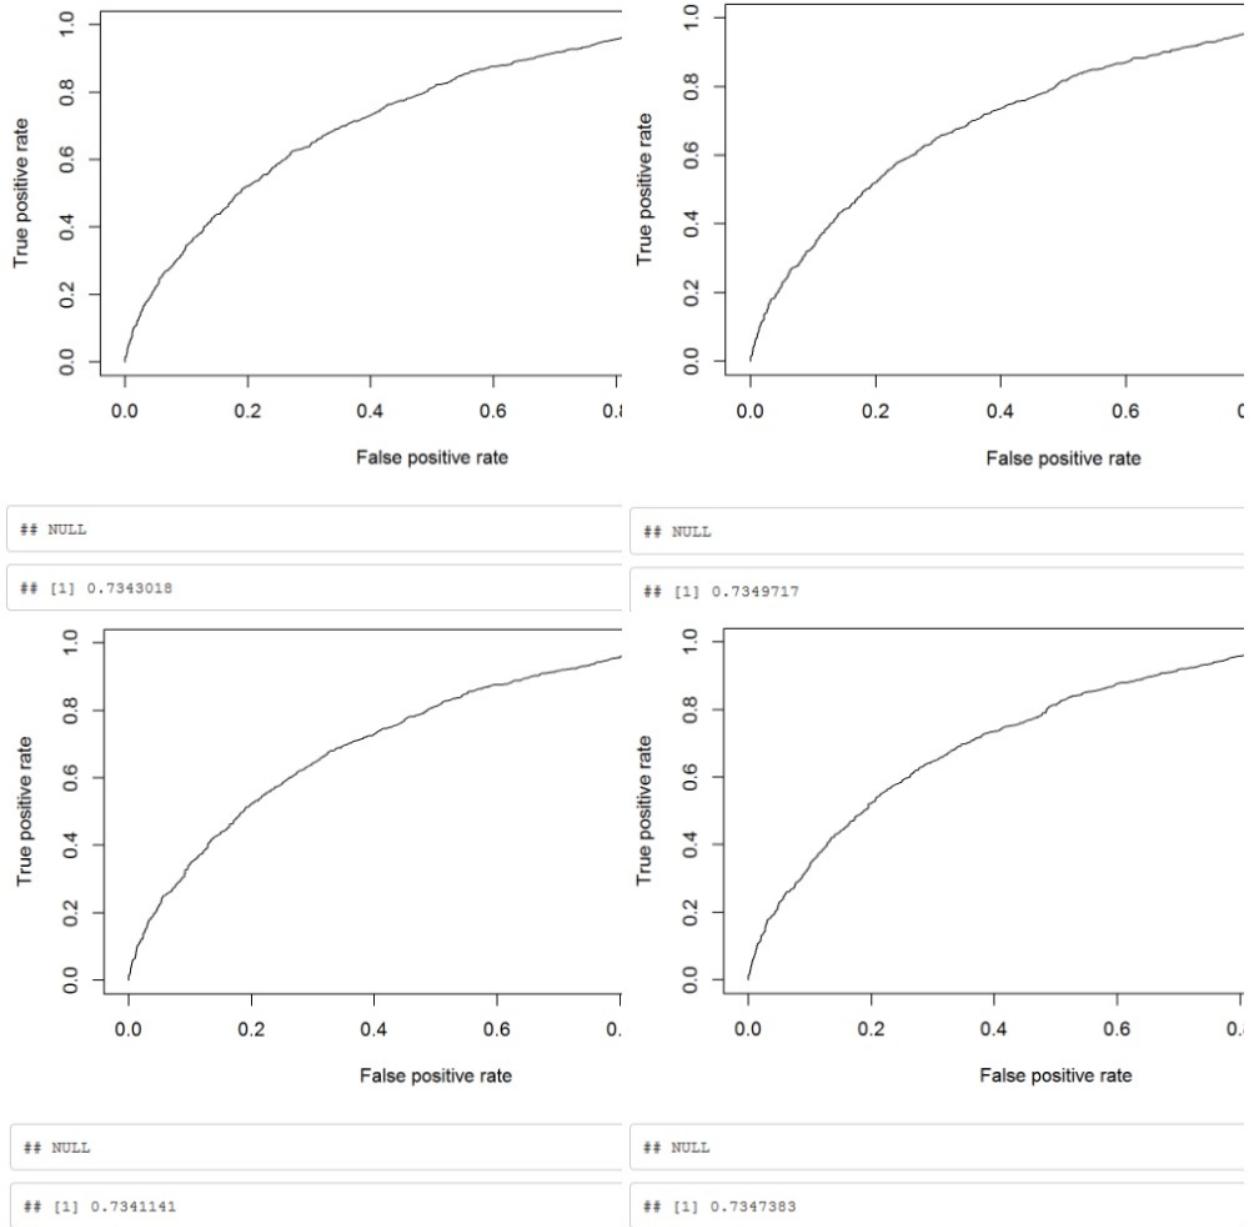
```
drops <- c("REVOKE")
ins.train <- ins.train[ , !(names(ins.train) %in% drops)]
ins.test <- ins.test[ , !(names(ins.test) %in% drops)]
```

So now that we have lowered our number of categorical variables, and thus lowered the total number of possible categorical combinations to calculate, we can use a tool “grind out” every combination and evaluate based on whatever criteria we wish:

Try out all categorical combinations with a tools, in this case - MuMIN DREDGE

(Intercept)	AGE	BLUEBOOK	CAR_AGE	CAR_TYPE	CLM_FREQ	HOME_VA	HOMEKIDS	INCOME	KIDSDRV	MSTATUS	MVR_PTS	OLDCLAIM	TIF	TRAVTIME	Y0J	R^2	adjR^2	df	logLik	AIC	delta	weight
0.2863497	-0.0021190	-1.7e-06	-0.0032151	0.0747763	0.0544899	-1e-07	0.0151576	-8.0e-07	0.0474560	0.1034545	0.0280908	NA	-0.006411	0.0013458	NA	0.1365173	0.1953955	15	-2227.142	4484.285	0.0000000	0.1493851
0.2769883	-0.0022200	-1.7e-06	-0.0031492	0.0748476	0.0548672	NA	0.0153841	-1.0e-06	0.0476973	0.1177475	0.0281612	NA	-0.006459	0.0013515	NA	0.1360284	0.1946957	14	-2228.340	4484.680	0.3950048	0.1226120
0.2952994	-0.0019948	-1.7e-06	-0.0033154	0.0746645	0.0544005	-1e-07	0.0163791	-7.0e-07	0.0469950	0.1022079	0.0280951	NA	-0.0065679	0.0013436	-0.0016678	0.1367252	0.1956930	16	-2226.633	4485.266	0.9813406	0.0914559
0.2869825	-0.0020824	-1.7e-06	-0.0032598	0.0747244	0.0547569	NA	0.0166928	-9.0e-07	0.0471914	0.1158737	0.0281633	NA	-0.0065669	0.0013490	-0.0017978	0.1362706	0.1950424	15	-2227.747	4485.493	1.2087644	0.0816258
0.2854805	-0.0021224	-1.7e-06	-0.0032063	0.0749379	0.0552478	-1e-07	0.0151184	-8.0e-07	0.0475379	0.1035760	0.0278606	6e-07	-0.0063310	0.0013572	NA	0.1366180	0.1955393	16	-2226.896	4485.791	1.5068840	0.0703220
0.2761348	-0.0022233	-1.7e-06	-0.0031405	0.0750113	0.0527742	NA	0.0153436	-1.0e-06	0.0477796	0.1178275	0.0279275	6e-07	-0.0063558	0.0013631	NA	0.1361320	0.1948440	15	-2228.086	4486.172	1.8877955	0.0581270
0.2718426	-0.0022592	NA	-0.0031631	0.0741195	0.0548951	-1e-07	0.0152289	-9.0e-07	0.0473281	0.1044133	0.0282627	NA	-0.0065659	0.0013556	NA	0.1356661	0.1941772	14	-2229.227	4486.453	2.1686555	0.0505114
0.2944955	-0.0019971	-1.7e-06	-0.0033074	0.0748281	0.0522992	-1e-07	0.0163502	-7.0e-07	0.0470742	0.1023203	0.02878607	6e-07	-0.0065570	0.0013552	-0.0016830	0.1368296	0.1958425	17	-2226.377	4486.754	2.4696220	0.0434546
0.2627404	-0.0023576	NA	-0.0030986	0.0741953	0.0552634	NA	0.0154515	-1.1e-06	0.0475669	0.1182205	0.0283307	NA	-0.0066616	0.0013611	NA	0.1351916	0.1934980	13	-2230.388	4486.776	2.4912182	0.0429879
0.3333078	-0.0029067	-1.7e-06	-0.0033371	0.0745735	0.0547403	-1e-07	NA	-8.0e-07	0.0614382	0.1009911	0.0281044	NA	-0.0066282	0.0013406	NA	0.1355962	0.1940772	14	-2229.398	4486.795	2.5109437	0.0425660
0.2861958	-0.0020844	-1.7e-06	-0.0032518	0.0748903	0.0526226	NA	0.0166620	-9.0e-07	0.0472711	0.1159396	0.0279251	6e-07	-0.0065559	0.0013607	-0.0018128	0.1363782	0.1951965	16	-2227.483	4486.966	2.6817581	0.0390815
0.3244006	-0.0030227	-1.7e-06	-0.0032712	0.0746438	0.0551322	NA	NA	-1.0e-06	0.0619013	0.1156548	0.0281771	NA	-0.0066330	0.0013464	NA	0.1350792	0.1933371	13	-2230.663	4487.326	3.0410358	0.0326554
0.2807023	-0.0021362	NA	-0.0032626	0.0740074	0.0548070	-1e-07	0.0164416	-8.0e-07	0.0468702	0.1029069	0.0282673	NA	-0.0065843	0.0013534	-0.0016558	0.1358710	0.1944708	15	-2228.725	4487.450	3.1655333	0.0306831
0.2727265	-0.0022213	NA	-0.0032082	0.0740717	0.0551548	NA	0.0167497	-1.0e-06	0.0470647	0.1163622	0.0283331	NA	-0.0065832	0.0013586	-0.0017839	0.1354300	0.1938393	14	-2229.804	4487.609	3.3243604	0.0283421
0.2709829	-0.0022625	NA	-0.0031542	0.0742638	0.0528099	-1e-07	0.0151891	-9.0e-07	0.0474111	0.1042652	0.0280297	6e-07	-0.0066468	0.0013671	NA	0.1357690	0.1943245	15	-2228.975	4487.950	3.6651398	0.0239019
0.2618968	-0.0023606	NA	-0.0030899	0.0743617	0.0531479	NA	0.0154105	-1.1e-06	0.0476503	0.1183007	0.0280942	6e-07	-0.0065513	0.0013728	NA	0.1352973	0.1936493	14	-2230.129	4488.258	3.9737720	0.0204839
0.3322936	-0.0029081	-1.7e-06	-0.0033278	0.0747396	0.0526288	-1e-07	NA	-8.0e-07	0.0614850	0.1011220	0.0278687	6e-07	-0.0066179	0.0013523	NA	0.1357017	0.1942281	15	-2229.140	4488.279	3.9947610	0.0202701
0.3412656	-0.0028688	-1.7e-06	-0.0034060	0.0744935	0.0546970	-1e-07	NA	-7.0e-07	0.0618542	0.1000879	0.0281078	NA	-0.0065818	0.0013390	-0.0010423	0.1356797	0.1941967	15	-2229.193	4488.387	4.1021235	0.0192107
0.1942740	NA	-1.9e-06	-0.0035167	0.0742477	0.0540864	-1e-07	0.0231384	-8.0e-07	0.0416199	0.1057101	0.0284068	NA	-0.0066501	0.0013247	NA	0.1352127	0.1935283	14	-2230.336	4488.672	4.3876610	0.0166548
0.3233988	-0.0030238	-1.7e-06	-0.0032620	0.0748120	0.0529882	NA	NA	-1.0e-06	0.0619474	0.1157424	0.0279376	6e-07	-0.0066226	0.0013582	NA	0.1351878	0.1934926	14	-2230.397	4488.794	4.5096734	0.0156691

The top 4 Models returned:



All in all, this approach yielded models with AUC's around 0.735, slightly higher than the basic step reduction.

4.2 TARGET_AMT

This is the target field that says whether the customer had to pay some amount after an accident. This field will only have a value if the TARGET_FLAG field has a 1. If the TARGET_FLAG field is a 0, then this field will be 0 as well. The first thing that we have to do is re-pick the training set. We do not want to use the exact same training set as before, because it is the same data and we really are not changing anything from the first models. We will be using a 70/30 split just like before.

4.2.1 Model 1

The first model that will be used is a form of a stepwise function. Basically, we select a set of fields from the training set and use that as a base model. We take a look at that model and see what fields should be kept and what fields should be removed. We remove the fields that we feel are not well correlated and get the final model.

The output for the model is as follows:

```
##  
## Call:  
## lm(formula = TARGET_AMT ~ PARENT1 + HOME_VAL + MSTATUS + TRAVTIME +  
##      CAR_USE + REVOKED + CAR_AGE + URBANICITY, data = training_2a)  
##  
## Residuals:  
##    Min     1Q Median     3Q    Max  
## -4448  -1728   -882    276  82773  
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)  
## (Intercept) 1.513e+03 3.582e+02  4.225 2.44e-05 ***  
## PARENT1     -8.084e+02 2.351e+02 -3.439 0.000589 ***  
## HOME_VAL    -3.020e-03 6.669e-04 -4.528 6.11e-06 ***  
## MSTATUS      1.967e+02 1.811e+02  1.086 0.277363  
## TRAVTIME     1.791e+01 4.432e+00  4.041 5.43e-05 ***  
## CAR_USE      -8.532e+02 1.496e+02 -5.701 1.27e-08 ***  
## REVOKED      5.913e+02 2.149e+02  2.751 0.005961 **  
## CAR_AGE      -5.571e+01 1.325e+01 -4.204 2.68e-05 ***  
## URBANICITY    1.809e+03 1.757e+02 10.293 < 2e-16 ***  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 4544 on 4222 degrees of freedom  
## Multiple R-squared:  0.05564,    Adjusted R-squared:  0.05385  
## F-statistic:  31.1 on 8 and 4222 DF,  p-value: < 2.2e-16
```

4.2.2 Model 2

The second model that we will be using is a forward selection, like used in the TARGET_FLAG prediction section above. This model takes a “blank” equation and starts to add variables until it finds the optimal solution for the model. It is a very similar process to the first model.

The output for the model is as follows:

```
##  
## Call:  
## lm(formula = TARGET_AMT ~ MVR PTS + URBANICITY + HOME_VAL + CAR_USE +  
##      CAR_AGE + PARENT1 + TRAVTIME + REVOKED + JOB + CAR_TYPE +  
##      TIF + CLM_FREQ, data = train2)  
##  
## Residuals:  
##    Min     1Q Median     3Q    Max  
## -6032  -1680   -811    277  83190  
##
```

```

## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) 1.407e+03 3.355e+02 4.194 2.79e-05 ***
## MVR PTS     1.721e+02 3.551e+01 4.848 1.30e-06 ***
## URBANICITY  1.620e+03 1.833e+02 8.837 < 2e-16 ***
## HOME_VAL    -2.819e-03 6.044e-04 -4.664 3.19e-06 ***
## CAR_USE     -6.736e+02 1.575e+02 -4.276 1.94e-05 ***
## CAR AGE     -4.563e+01 1.355e+01 -3.368 0.000765 ***
## PARENT1     -8.446e+02 2.119e+02 -3.986 6.82e-05 ***
## TRAVTIME    1.690e+01 4.416e+00 3.826 0.000132 ***
## REVOKED      5.478e+02 2.142e+02 2.557 0.010583 *
## JOB          -2.734e+02 1.508e+02 -1.813 0.069946 .
## CAR_TYPE     2.641e+02 1.530e+02 1.727 0.084295 .
## TIF          -2.886e+01 1.678e+01 -1.720 0.085529 .
## CLM_FREQ     9.993e+01 6.786e+01 1.473 0.140924
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4522 on 4218 degrees of freedom
## Multiple R-squared: 0.06572, Adjusted R-squared: 0.06307
## F-statistic: 24.73 on 12 and 4218 DF, p-value: < 2.2e-16

```

It is very interesting to compare the two models in this section. We can see by the summary statistics of both models, that Model 1's coefficients are statistically significant (below .05 confidence), while model 1 seems to have a few variables that are not significant at all, but the R squared value for model 1 is lower than the R squared of the second model. It will be interesting to do some further analysis and see which model is actually better than the other.

5 Select Models

Now that all of the models have been created and predicted, it is time to pick and choose which are the best. We will pick the best model for TARGET_FLAG (probit model) and the best model for the TARGET_AMT field (linear regression).

5.1 TARGET_FLAG

5.1.1 F1 Score

```

##      Predicted
## Actual   0   1
##       0 3113  13
##       1  616 489

##      Predicted
## Actual   0   1
##       0 2907 219
##       1  647 458

##      Predicted
## Actual   0   1
##       0 2907 219
##       1  647 458

```

We can take a look at how well the models compare to each other. One way to do that is to look at the F1 score. This score takes the precision and recall of the model and lets us know how the model fairs. The equation for the F1 score is as follows:

$$F1Score = \frac{2 * precision * recall}{precision + recall}$$

We can see that all of the models have pretty high F1 scores. The best most is model 1 by just a little bit.

model	precision	recall	F1 Score
model 1	.8563123	.9974202	.9214956
model 2	.8722322	.9316075	.9009427
model 3	.8722322	.9316075	.9009427

5.1.2 ROC Curve

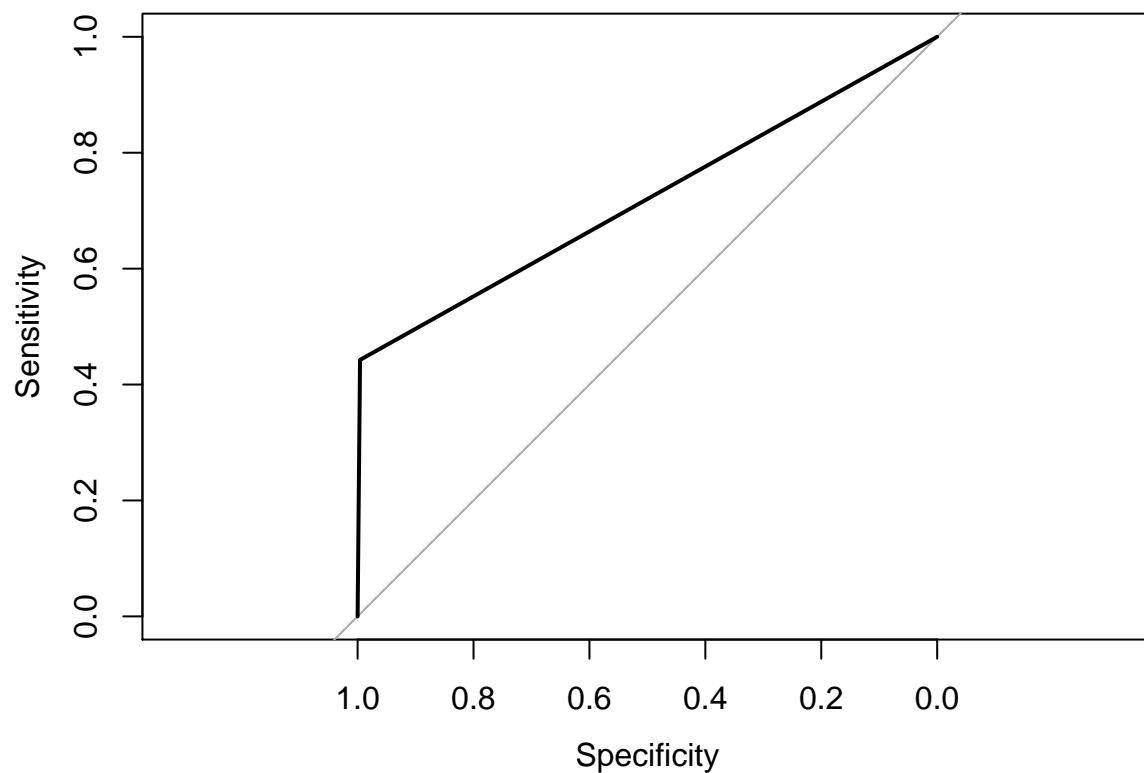
Another way to look at the models is through the ROC curve. The ROC curve compares the sensitivity of the model with the specificity of the model. It basically give the performance of the model. With that curve, we can calculate the AUC (Area Under the Curve). The higher this number is, the better the performance of the model. We can see that model 1 is still the best choice from all three models.

```
train3 <- cbind(train1 , answer1a[1:4231] , answer , answer2)

rc1 <- roc(factor(TARGET_FLAG) ~ answer1a[1:4231] , data=train3)
rc2 <- roc(factor(TARGET_FLAG) ~ answer , data=train3)
rc3 <- roc(factor(TARGET_FLAG) ~ answer2 , data=train3)

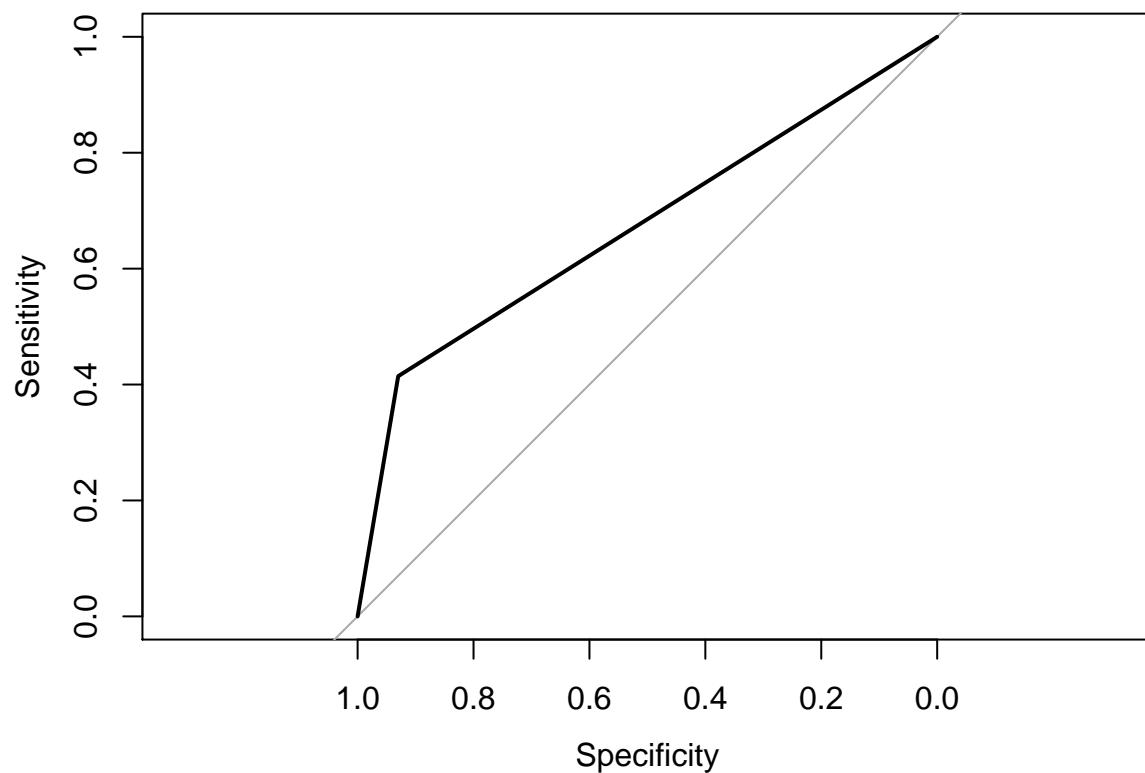
plot(rc1,main='Model 1 - ROC Curve')
```

Model 1 – ROC Curve



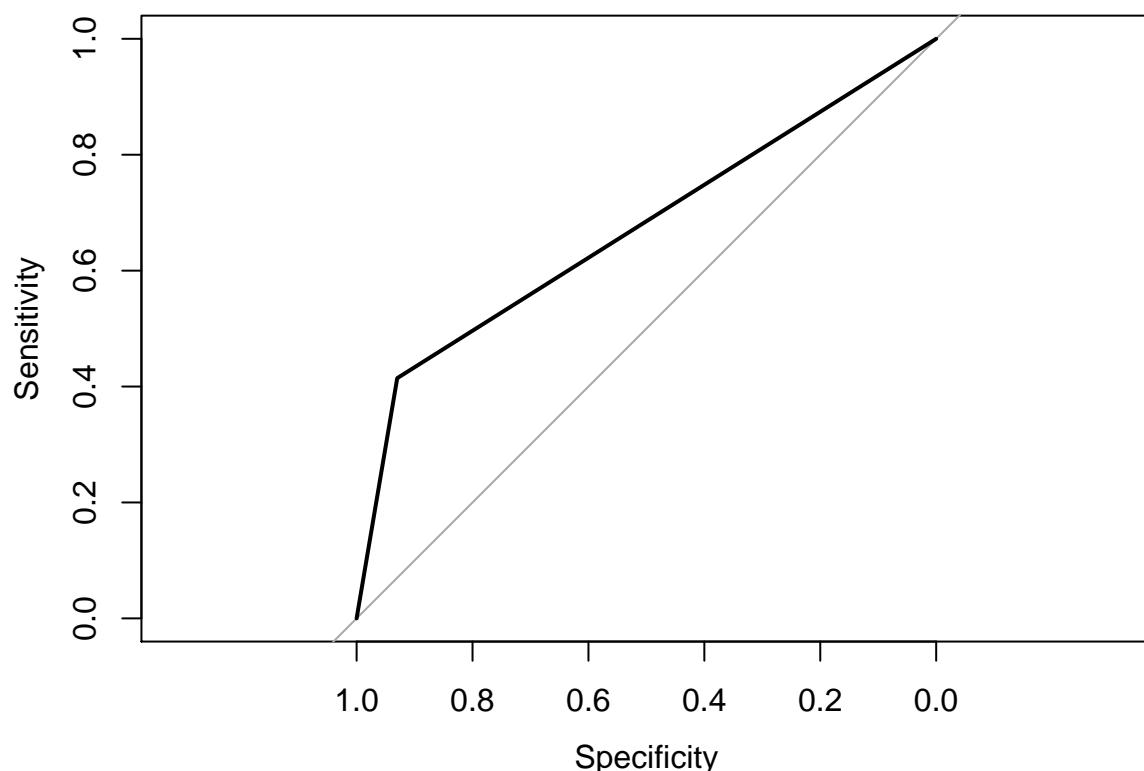
```
##  
## Call:  
## roc.formula(formula = factor(TARGET_FLAG) ~ answer1a[1:4231],      data = train3)  
##  
## Data: answer1a[1:4231] in 3126 controls (factor(TARGET_FLAG) 0) < 1105 cases (factor(TARGET_FLAG) 1)  
## Area under the curve: 0.7192  
  
plot(rc2,main='Model 2 - ROC Curve')
```

Model 2 – ROC Curve



```
##  
## Call:  
## roc.formula(formula = factor(TARGET_FLAG) ~ answer, data = train3)  
##  
## Data: answer in 3126 controls (factor(TARGET_FLAG) 0) < 1105 cases (factor(TARGET_FLAG) 1).  
## Area under the curve: 0.6722  
  
plot(rc3,main='Model 3 - ROC Curve')
```

Model 3 – ROC Curve



```
##  
## Call:  
## roc.formula(formula = factor(TARGET_FLAG) ~ answer2, data = train3)  
##  
## Data: answer2 in 3126 controls (factor(TARGET_FLAG) 0) < 1105 cases (factor(TARGET_FLAG) 1).  
## Area under the curve: 0.6722  
  
model <- c('Model 1', 'Model 2', 'Model 3')  
area <- c(auc(train1$TARGET_FLAG, answer1a), auc(train1$TARGET_FLAG, answer), auc(train1$TARGET_FLAG, ans  
df <- data.frame(Model=model, AUC=area)  
df  
  
##      Model      AUC  
## 1 Model 1 0.7191876  
## 2 Model 2 0.6722110  
## 3 Model 3 0.6722110  
  
model | AUC (Area Under Curve)  
-----|-----  
model 1 | .7283561  
model 2 | .6572123  
model 3 | .6572123
```

5.1.3 AIC/BIC/Log-Likelihood

The Akaike information criterion (AIC) is a measure of the relative quality of statistical models for a given set of data. Given a collection of models for the data, AIC estimates the quality of each model, relative to

each of the other models. Hence, AIC provides a means for model selection. The lower the AIC the better the model.

Bayesian information criterion (BIC), is a criterion for model selection among a finite set of models. The lower the BIC, the better the model. Ghand and hand with AIC.

We can see that the smallest AIC and BIC is still model number 1. All of the criteria is pointing towards model number 1. That means that the best model, out of these three would be model number 1 to continue on for the evaluation set.

model	AIC	BIC
model 1	3188.62	3258.47
model 2	2900.54	4021.20
model 3	82733.65	82835.25

5.2 TARGET_AMT

We first check the summary stats with the two models. The first this we check is the MSE (Mean Squared Error). This is the mean of the residuals (actual - predicted) squared. It is a good way to see how accurate your model is. A smaller MSE is always good. The next things is the R squared. This is usually called the goodness of fit. The higher the R squared value the better the model is. The last thing is the F-Stat. It is most often used when comparing statistical models that have been fitted to a data set, in order to identify the model that best fits the population from which the data were sampled.

All three of these parameters are showing that the best model to use would be model 1. It has the lowest MSE, the highest R squared and the highest F-stat (which means a lower alpha value and most statistically relevant).

model	MSE (Mean Squared Error)	R Squared	F-Stat
model 1	18051149	.06530975	86.13564
model 2	18329321	.05090597	40.27218

Now we can look at the residuals and see how well they fit to the regression line. In figure 4.1, we look at the first model. Overall, the data does not fit a linear model at all. The histograms of the residual (lower left have corner) are heavily skewed to the right. Almost all of the data is clustered around 0. Next, the normal plot (top left corner) has a pretty obvious bend to it. The line is supposed to be pretty straing. Finally the residual plot (top right corner) show a pretty obvious patter amoungst the points. It is supposed to be random. This is showing that the data is really not lending itself to a linear model. That means that a different model may be a better choice. In figure 4.2, we see the exact same patterns as in Figure 4.1. This data does not lend itself to a linear model.

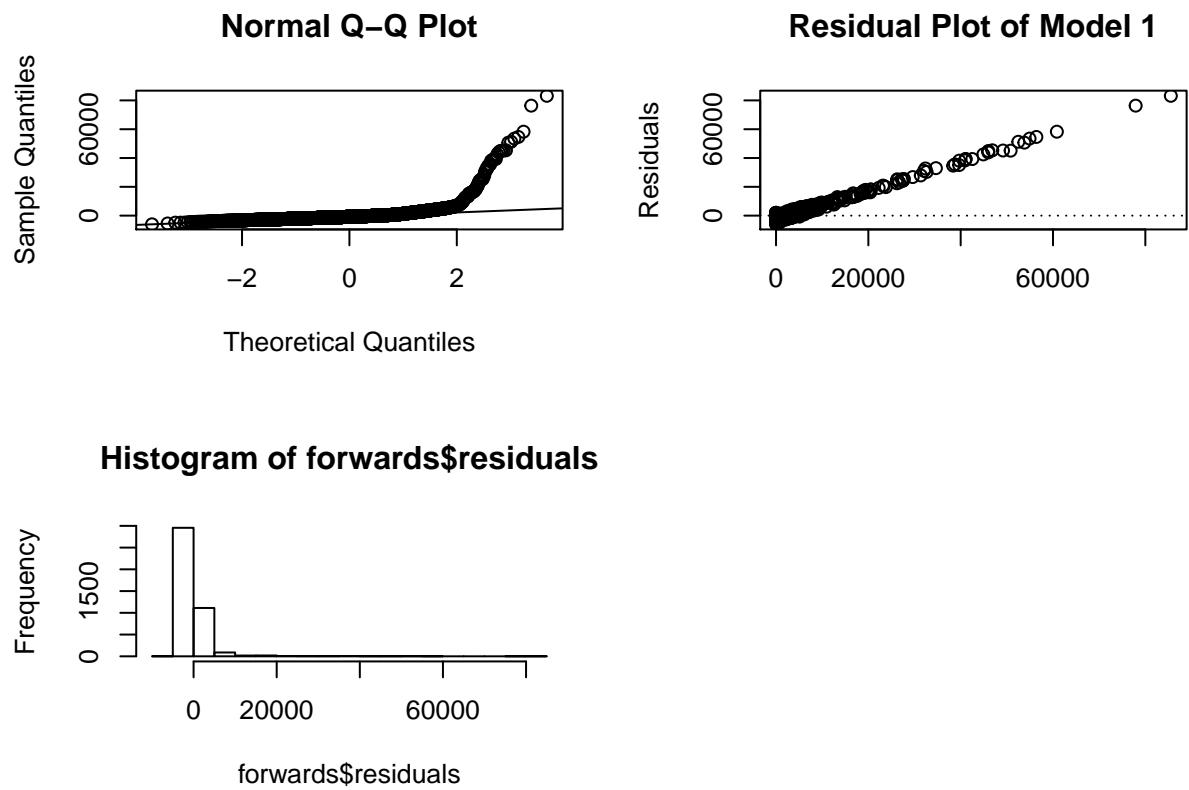


Figure 4.1

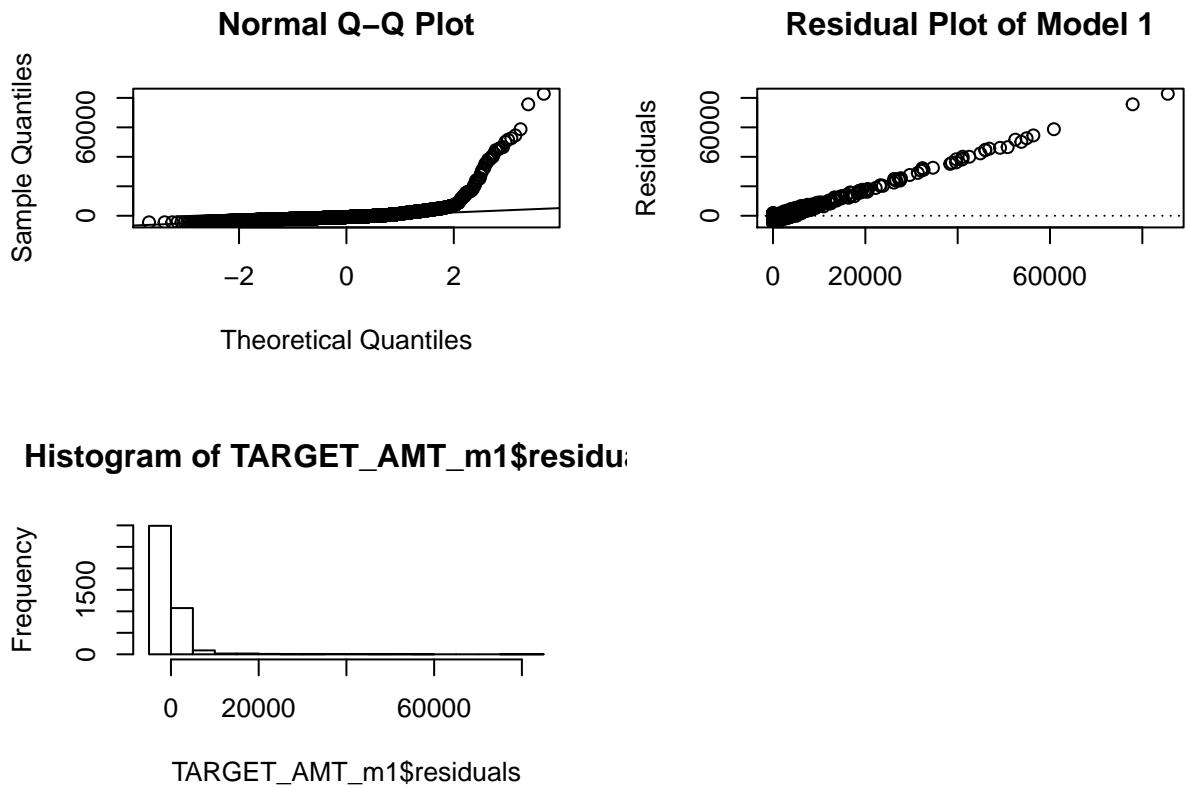


Figure 4.2

Overall, model 1 is the best choice. The data may not lend itself to a linear model, but it is a better choice than Model 2.

That means, to predict the TARGET_FLAG, we will be using model 1 and to predict the TARGET_AMT we will also be using model 1.

6 Model Evaluation

The final flag values and the amounts can be seen in the final_result.csv

```
## final_answer
##    0    1
## 1696   19
```

7 Code Appendix

```
library(PerformanceAnalytics)
library(ggplot2)
library(gridExtra)
library(knitr)
library(lattice)
library(tidyr)
```

```

library(dplyr)
library(pROC)
library(stringr)
library(aod)
library(Rcpp)
library(Amelia)
library(corrplot)
library(gam)
library(pscl)
library(ROCR)
library(gmodels)
library(rpart)
library(Metrics)

data <- read.csv('https://raw.githubusercontent.com/jhamski/DATA621-Homework/master/Homework_4/insurance.csv')

missmap(data, legend = TRUE, main = "Missing Values vs Observed", col = c("white", "black"))

summary(data)

summary(data$TARGET_FLAG)
hist(data$TARGET_FLAG)

summary(data$TARGET_AMT)
hist(data$TARGET_AMT)

blue_book <- unname(sapply(data$BLUEBOOK, str_replace_all, '[,$]', ''))
blue_book <- as.numeric(blue_book)

income <- unname(sapply(data$INCOME, str_replace_all, '[,$]', ''))
income <- as.numeric(income)

home_val <- unname(sapply(data$HOME_VAL, str_replace_all, '[,$]', ''))
home_val <- as.numeric(home_val)

old_claim <- unname(sapply(data$OLDCLAIM, str_replace_all, '[,$]', ''))
old_claim <- as.numeric(old_claim)

data$BLUEBOOK <- blue_book
data$INCOME <- income
data$HOME_VAL <- home_val
data$OLDCLAIM <- old_claim

data2 <- data[,-c(1,2,3,9,11,12,13,14,16,19,20,23,26)]
chart.Correlation(data2)

the_cor <- cor(data[sapply(data, is.numeric)])
#the_cor
corrplot(the_cor, method = "circle")

data <- data[complete.cases(data),]
data <- data[data$CAR_AGE >= 0,]

```

```

data$PARENT1 <- ifelse(data$PARENT1 == "No", 1, 0)
data$SEX <- ifelse(data$SEX == 'M', 0, 1)
data$CAR_USE <- ifelse(data$CAR_USE == 'Commercial', 0, 1)
data$MSTATUS <- ifelse(data$MSTATUS == 'Yes', 0, 1)
data$RED_CAR <- ifelse(data$RED_CAR == "no", 0, 1)
data$EDUCATION <- ifelse(data$EDUCATION %in% c('PhD', "Masters"), 0, 1)
data$REVOKE <- ifelse(data$REVOKE == "No", 0, 1)
data$URBANICITY <- ifelse(data$URBANICITY == "Highly Urban/ Urban", 1, 0)
data$JOB <- ifelse(data$JOB %in% c('Professional', 'Manager', 'Student', 'Lawyer'), 1, 0)
data$CAR_TYPE <- ifelse(data$CAR_TYPE %in% c('Panel Truck', "Pickup", "Sports Car"), 1, 0)

data <- data[,-1]
data <- data[sample(nrow(data)),]
top <- round(.70 * NROW(data))

train1 <- data[1:top,]
test1 <- data[(top + 1):NROW(data),]

training_2a <- dplyr::select(train1, -c(KIDSDRV, HOMEKIDS, EDUCATION, JOB, TIF,
                                         CAR_TYPE, OLDCLAIM, CLM_FREQ, MVR PTS))
M11 <- lm(TARGET_FLAG ~ . - TARGET_FLAG, data = training_2a)
#summary(M11)
M12 <- update(M11, . ~ . - AGE - HOMEKIDS - YOJ - INCOME - SEX - EDUCATION - BLUEBOOK - RED_CAR - OLDCLAIM - CLM_FREQ)
#summary(M12)
TARGET_FLAG_m1 <- M12
summary(TARGET_FLAG_m1)

answer1a <- predict(TARGET_FLAG_m1, type = "response")
answer1a <- ifelse(answer1a < .5, 0, 1)

fullmod <- glm(TARGET_FLAG ~ KIDSDRV + AGE + HOMEKIDS + YOJ + INCOME + PARENT1 + HOME_VAL + MSTATUS + +
                 CAR_TYPE + OLDCLAIM + CLM_FREQ + MVR PTS)

backwards <- step(fullmod, trace = 0)
prediction <- round(predict(backwards, type = 'response'), 4)

answer <- ifelse(prediction < .5, 0, 1)

summary(backwards)

nothing <- glm(TARGET_FLAG ~ 1, data = train1, family = binomial(link = 'probit'))
forwards <- step(nothing, scope = list(lower = formula(nothing), upper = formula(fullmod)), direction = "for")

pred <- round(predict(forwards, type = 'response'), 4)

answer2 <- ifelse(pred < .5, 0, 1)

summary(forwards)

data <- data[sample(nrow(data)),]
top <- round(.70 * NROW(data))

train2 <- data[1:top,]
test2 <- data[(top + 1):NROW(data),]

```

```

training_2a <- dplyr::select(train2, -c(KIDSDRV, HOMEKIDS, EDUCATION, JOB, TIF,
                                         CAR_TYPE, OLDCLAIM, CLM_FREQ, MVR PTS))
M11 <- lm(TARGET_AMT ~ . - TARGET_FLAG, data=training_2a)
M12 <- update(M11, . ~ . - AGE - HOMEKIDS - YOJ - INCOME - SEX - EDUCATION - BLUEBOOK - RED_CAR - OLDCLAIM - CLM_FREQ)
TARGET_AMT_m1 <- M12

pred5 <- predict(TARGET_AMT_m1)
summary(TARGET_AMT_m1)

nothing <- lm(TARGET_AMT ~ 1, data = train2)
forwards <- step(nothing, scope = list(lower=formula(nothing), upper=formula(fullmod)), direction = "forwards")

pred4 <- predict(forwards)
summary(forwards)

t1 <- table(Actual = train1$TARGET_FLAG, Predicted = answer1a)
t1
t2 <- table(Actual = train1$TARGET_FLAG, Predicted = answer)
t2
t3 <- table(Actual = train1$TARGET_FLAG, Predicted = answer2)
t3

fscore <- function(x)
{
  precision <- x[1]/(x[1] + x[4])
  recall <- x[1] / (x[1] + x[3])
  score <- (2*precision*recall)/(precision+recall)

  return(c(precision, recall, score))
}

model1 <- fscore(t1)
model2 <- fscore(t2)
model3 <- fscore(t3)

train3 <- cbind(train1, answer1a[1:4231], answer, answer2)

rc1 <- roc(factor(TARGET_FLAG) ~ answer1a[1:4231], data=train3)
rc2 <- roc(factor(TARGET_FLAG) ~ answer, data=train3)
rc3 <- roc(factor(TARGET_FLAG) ~ answer2, data=train3)

plot(rc1, main='Model 1 - ROC Curve')
plot(rc2, main='Model 2 - ROC Curve')
plot(rc3, main='Model 3 - ROC Curve')

model <- c('Model 1', 'Model 2', 'Model 3')
area <- c(auc(train1$TARGET_FLAG, answer1a), auc(train1$TARGET_FLAG, answer), auc(train1$TARGET_FLAG, answer2))
df <- data.frame(Model=model, AUC=area)
df

# Akaike Information Criterion
AIC.1 <- AIC(TARGET_FLAG_m1)

```

```

AIC.2 <- AIC(backwards)
AIC.3 <- AIC(forwards)
AIC <- rbind(AIC.1, AIC.2, AIC.3) %>% round(2)

# BIC
BIC.1 <- BIC(TARGET_FLAG_m1)
BIC.2 <- BIC(backwards)
BIC.3 <- BIC(forwards)
BIC <- rbind(BIC.1, BIC.2, BIC.3) %>% round(2)

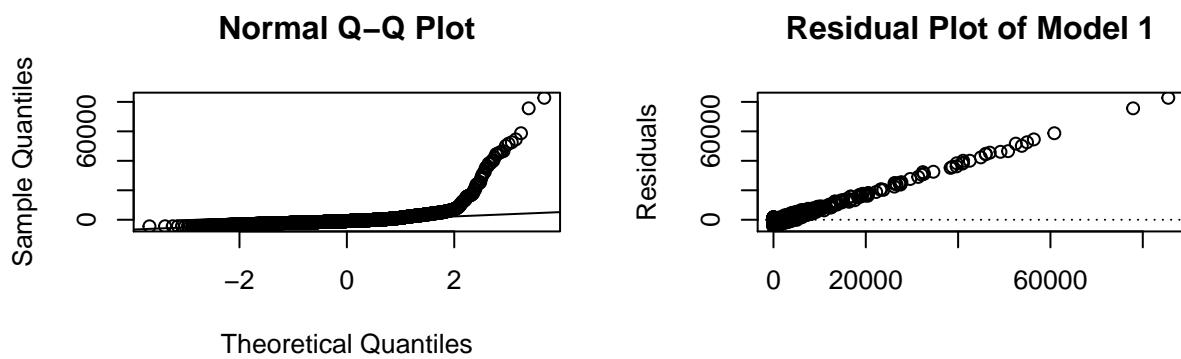
model1MSE <- Metrics::mse(train2$TARGET_AMT, pred4)
model2MSE <- Metrics::mse(train2$TARGET_AMT, pred5)

col_mse <- c(model1MSE,model2MSE)
col_r_sq <- c(summary(forwards)$r.squared, summary(TARGET_AMT_m1)$r.squared)
col_f_stat <- c(summary(aov(forwards))[[1]]$F[1], summary(aov(TARGET_AMT_m1))[[1]]$F[1])

par(mfrow = c(2,2))
g1<- qqnorm(forwards$residuals)
g2 <- qqline(forwards$residuals)
g3 <- plot(forwards$residuals ~ train2$TARGET_AMT,
           xlab=' ',
           ylab='Residuals',
           main='Residual Plot of Model 1')
abline(h=0,lty=3)
g4 <- hist(forwards$residuals)

## final_answer
##      0      1
## 1696   19

```



Histogram of TARGET_AMT_m1\$residu

